

Computer Vision 1

THOMAS MENSINK
MASTER AI
UNIVERSITY OF AMSTERDAM

LECTURES/THEORY

- 06-02-2018, 17:00-19:00, C0.05, **Introduction** (Szeliski 1)
- 13-02-2018, 17:00-19:00, C0.05, **Image Formation** (Szeliski: 2.I.1 + 2.I.2 + 2.2 + 2.3.2 + 2.3.3)
- 20-02-2018, 17:00-19:00, C0.05, **Color and Image Processing** (Szeliski: 3.I + 3.2 + 3.3)
- 27-02-2018, 17:00-19:00, C0.05, **Feature Detection, Motion and Classification** (Szeliski: 4, 8.I.1 + 8.I.3 + 8.2.I + 8.4; Bengio: 4 + 5.I + 5.2 + 5.3 + 5.7 + 5.8 + 5.9)
- 06-03-2018, 17:00-19:00, C0.05, **Object Recognition: BoW and ConvNets** (Szeliski: 5.I.1 + 5.I.4 + 5.I.5 + 5.2 + 5.3 + 5.4, 6.I + 6.3, 14.I + 14.2.I + 14.3 + 14.4.I; Bengio: 7.2 + 7.4 + 9.I + 9.2 + 9.3)
- 13-03-2018, 17:00-19:00, C0.05, **Deep Learning, Stereo and 3D Reconstruction** (Szeliski: 11.I + 11.2 + 11.3 + 11.4, 12.I + 12.2; Bengio: 12.I + 12.2)
- 20-03-2018, 17:00-19:00, C0.05, **Applications** (Szeliski: 12.6.2 + 12.6.3 + 12.2.4)
- 26-03-2018, Monday, 9:00-12:00, **Written Exam**

TODAY: FEATURE DETECTION, MOTION AND CLASSIFICATION

Focus

Szeliski:

- 4.1 Point & Patches
- 4.2 Edges
- 4.3 Lines
- 4.4 Additional reading / exercises
- 8.1 Translational alignment (-Fourier)
- 8.2.1 Parametric motion: Application Video Stabilisation
- 8.4 Optical Flow

If we need it ...

Bengio:

- 4 Numerical Computation
- 5 Machine Learning Basics

Note: Should've been covered in ML1

LAST WEEK

Class

- Color invariants
- Image processing
- Filtering and convolutions

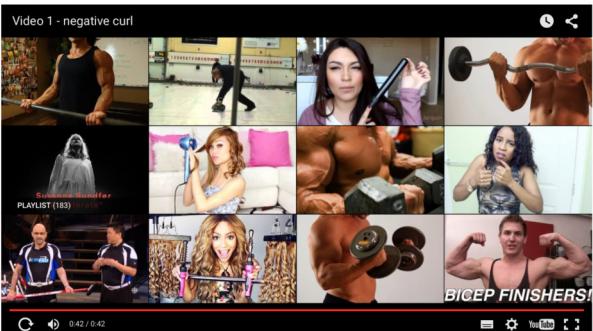
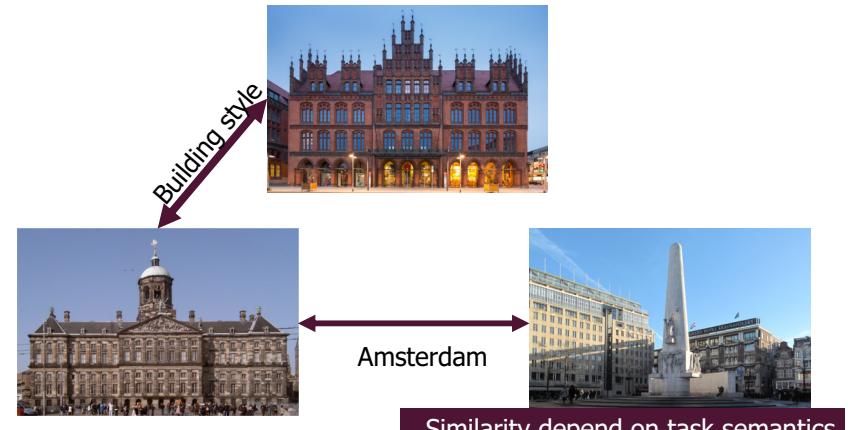
Lab

- Matlab assignment
- Piazza

TODAY'S CLASS

- Image Similarity
- Feature Detection
- Motion

IMAGE SIMILARITY



Source: NOS.nl | 21st Oct 2015

HOW TO REPRESENT AN IMAGE?

PIXEL VALUES



Rescale Image
64 x 64

173	69	201	220	215	106	64	194	47	142
234	40	232	105	225	18	168	148	223	92
1	173	49	189	14	91	130	199	106	52
12	233	86	222	75	21	155	216	158	40
156	182	192	191	237	177	180	7	122	58
65	170	185	3	158	112	230	116	19	213
180	88	80	90	255	64	163	177	194	9
79	47	97	80	250	106	210	81	45	211
2	236	52	126	170	184	155	120	240	244
196	219	237	147	123	14	232	216	59	25
246	218	217	82	228	87	239	58	95	157
224	246	204	229	80	76	95	92	42	36
121	53	6	44	78	239	41	3	22	219
183	206	236	125	135	22	12	103	131	86
56	184	63	80	230	187	240	94	254	166
217	240	242	129	91	247	49	37	217	175
39	24	225	179	163	65	35	29	215	246

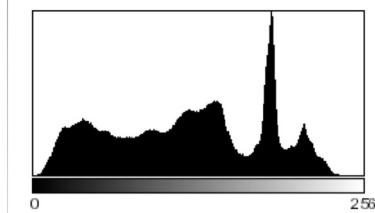
Concatenate to 12.888 dim vector

PIXEL VALUES - SIMILARITIES



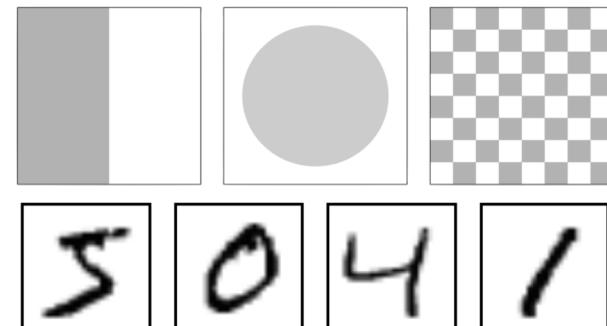
These three images have equal distance to (green) query image
Raw pixel values do not reflect visual similarity

GLOBAL HISTOGRAMS

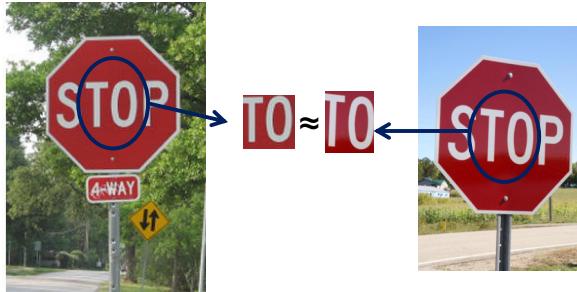


Count: 1920000
Mean: 118.848
StdDev: 59.179
Min: 0
Max: 251
Mode: 184 (30513)

GLOBAL HISTOGRAMS: EXAMPLES



LOCAL CORRESPONDENCE



Features are local, so robust to occlusion and clutter

TODAY'S CLASS

- Image Similarity
- Feature Detection
- Motion

FEATURE DETECTION

- Image alignment
- 3D reconstruction
- Motion tracking
- Robot navigation
- Indexing and database retrieval
- Object recognition



IMAGE MATCHING



by [Diva Sian](#)



by [swashford](#)

IMAGE MATCHING



by Diva Sian

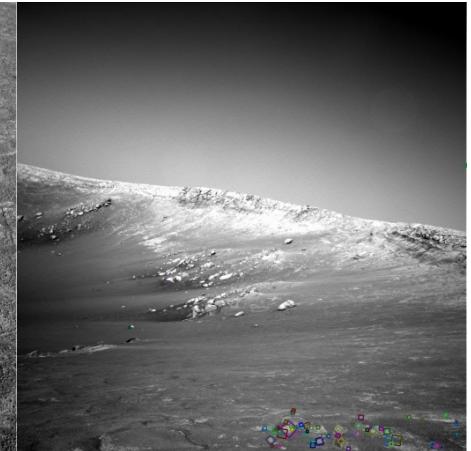


by scqbt

EVEN HARDER, STILL POSSIBLE?



NASA Mars Rover images



SIFT matches by Noah Snavely

Local features and alignment

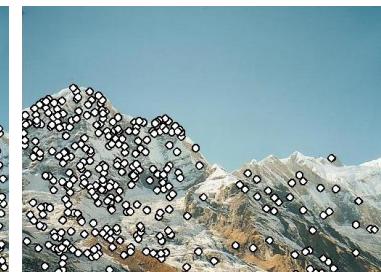


- We need to match (align) images
- Global methods sensitive to occlusion, lighting, parallax effects. So look for local features that match well.
- How would you do it by eye?

[Darya Frolova and Denis Simakov]

Local features and alignment

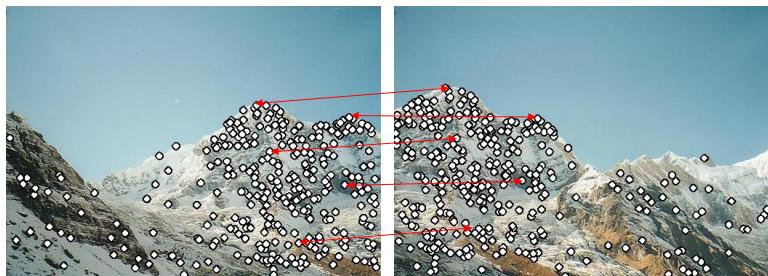
- Detect feature points in both images



[Darya Frolova and Denis Simakov]

Local features and alignment

- Detect feature points in both images
- Find corresponding pairs



[Darya Frolova and Denis Simakov]

Local features and alignment

- Detect feature points in both images
- Find corresponding pairs
- Use these pairs to align images



[Darya Frolova and Denis Simakov]

Local features and alignment

- Problem 1:
 - Detect the *same* point *independently* in both images



no chance to match!

We need a **repeatable** detector

[Darya Frolova and Denis Simakov]

Local features and alignment

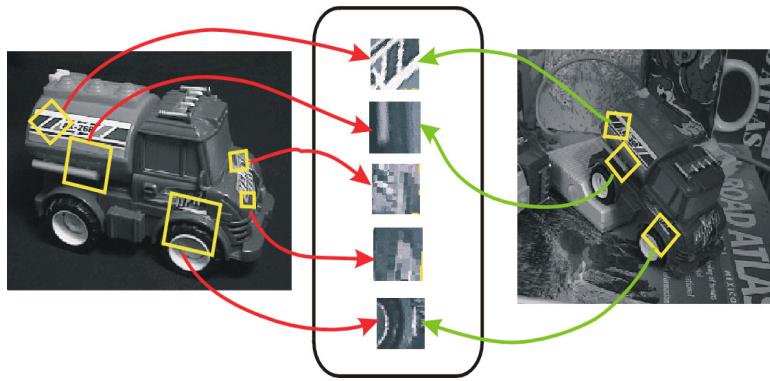
- Problem 2:
 - For each point correctly recognize the corresponding one



We need a **reliable** descriptor

[Darya Frolova and Denis Simakov]

INVARIANT LOCAL DESCRIPTORS



Local descriptors should transform image content to feature space:
invariant to translation, rotation, scale, and other imaging parameters

MAIN QUESTIONS

1. Where will the interest points come from?
2. How to describe a local region?
3. How to establish correspondences, i.e., compute matches?

Note: *interest points* are *key-points* (and sometimes called *features*)

Many applications

- tracking: which points are good to track?
- recognition: find patches likely to tell us something about object category
- 3D reconstruction: find correspondences across different views

Local features and alignment

- Problem 3:
 - Detect distinctive points in both images



no chance to find matching pairs!

We need a **distinctive** detector/**descriptor**

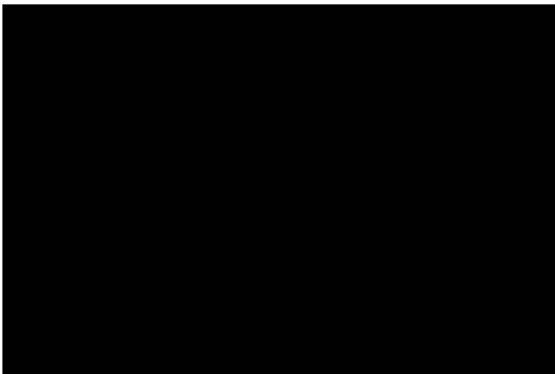
[Darya Frolova and Denis Simakov]

EXAMPLE: STRUCTURE FROM MOTION



<https://www.youtube.com/watch?v=4cEQZreQ2zQ>

EXAMPLE: STRUCTURE FROM MOTION



<https://www.youtube.com/watch?v=4cEOZreQ2zQ>

INTEREST POINTS



0D structure: **single points**
→ not useful for matching



1D structure: **lines**
→ edge, can be localised in 1D,
subject to the aperture problem



2D structure: **corners, blobs**
→ corner, or **interest point**, can be
localised in 2D, good for matching

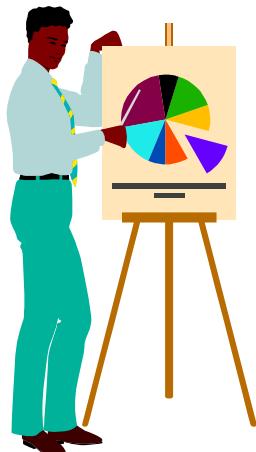
Interest Points have **2D** structure.

Where will interest points come from?

INTEREST POINT DETECTORS

1. Edges
2. Corners
3. Blobs
4. Templates

Canny Edge Detector



CANNY EDGE DETECTOR

- This is probably the most widely used edge detector in computer vision
- Theoretical model: step-edges corrupted by additive Gaussian noise
- Canny has shown that the first derivative of the Gaussian closely approximates the operator that optimizes the product of *signal-to-noise ratio* and localization

J. Canny, [A Computational Approach To Edge Detection](#), IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

DESIGNING AN EDGE DETECTOR

Criteria for a good edge detector:

- Good detection
 - the optimal detector should find all real edges, ignoring noise or other artifacts
- Good localization
 - the edges detected must be as close as possible to the true edges
 - the detector must return one point only for each true edge point

Cues of edge detection

- Differences in color, intensity, or texture across the boundary
- Continuity and closure
- High-level knowledge

Source: L. Fei-Fei

PIPELINE

1. Filter image with x, y derivatives of Gaussian
 2. Find magnitude and orientation of gradient
 3. Non-maximum suppression:
 - Thin multi-pixel wide “ridges” down to single pixel width
 4. Thresholding and linking (hysteresis):
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
- MATLAB: `edge(image, 'canny')`

Source: D. Lowe, L. Fei-Fei

Source: L. Fei-Fei

FINDING CANNY EDGES



COMPUTE GRADIENTS

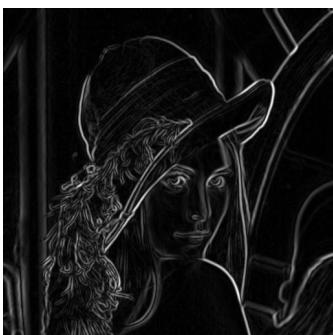


X-Derivative of Gaussian

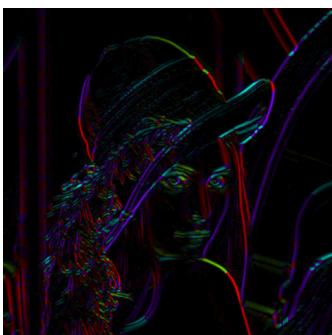


Y-Derivative of Gaussian

GET MAGNITUDE AND ORIENTATION

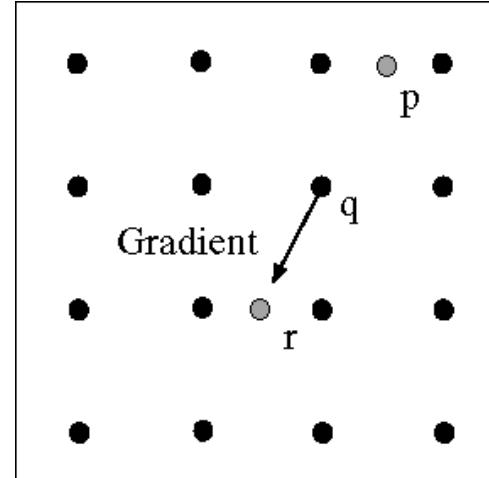


Gradient Magnitude

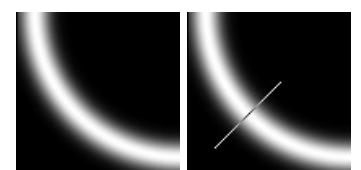


$\theta = \text{atan2}(g_y, g_x)$

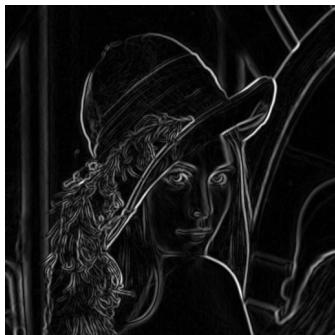
NON-MAXIMUM SUPPRESSION



At q, we have a maximum if the value is larger than those at both p and at r. Interpolate to get these values.



NON-MAXIMUM SUPPRESSION



Before



After

THRESHOLDING AND LINKING

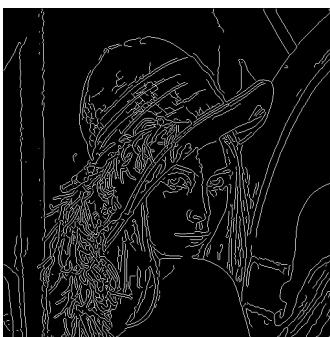


Before



After

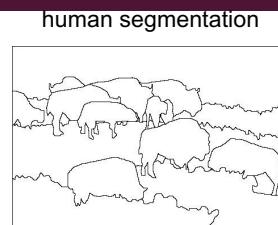
FINAL CANNY EDGES



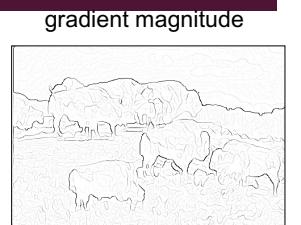
WHERE DO HUMANS SEE BOUNDARIES?



image



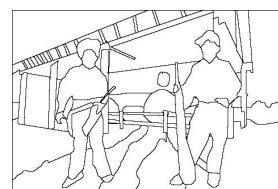
human segmentation

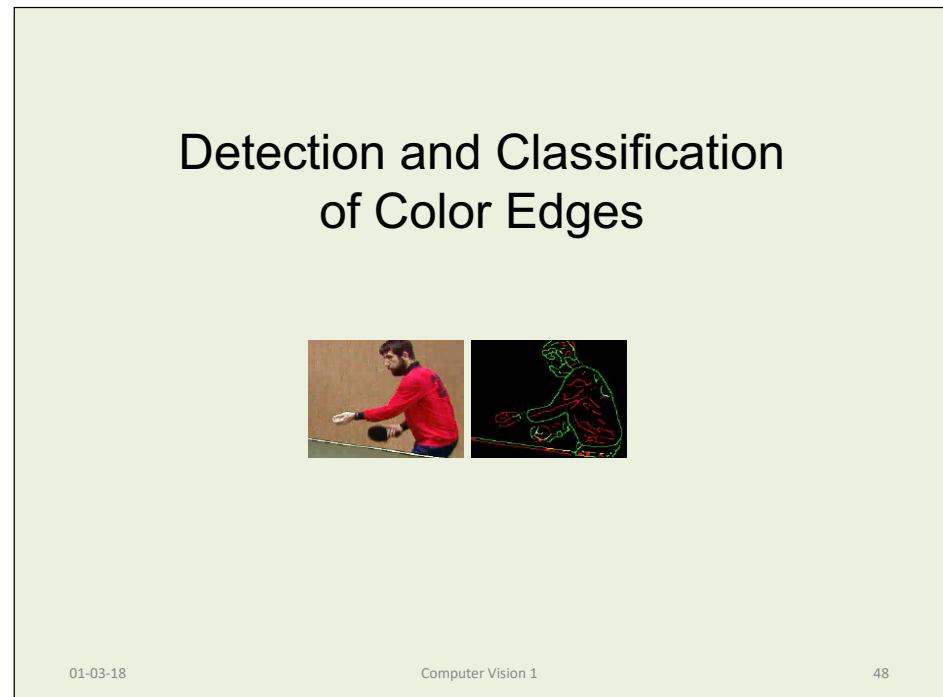
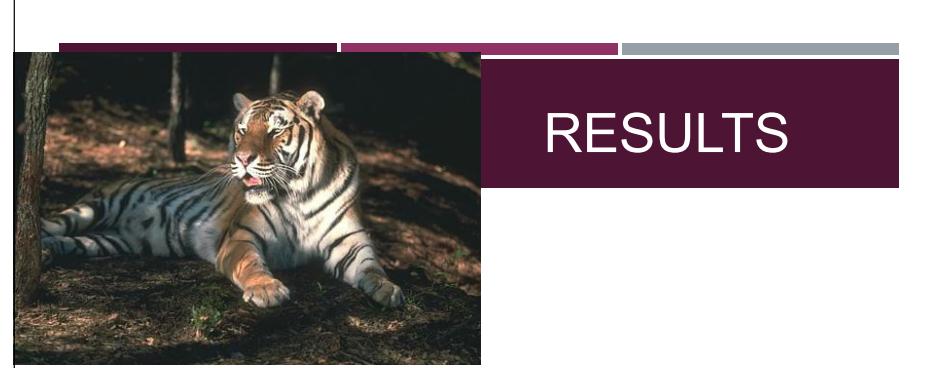
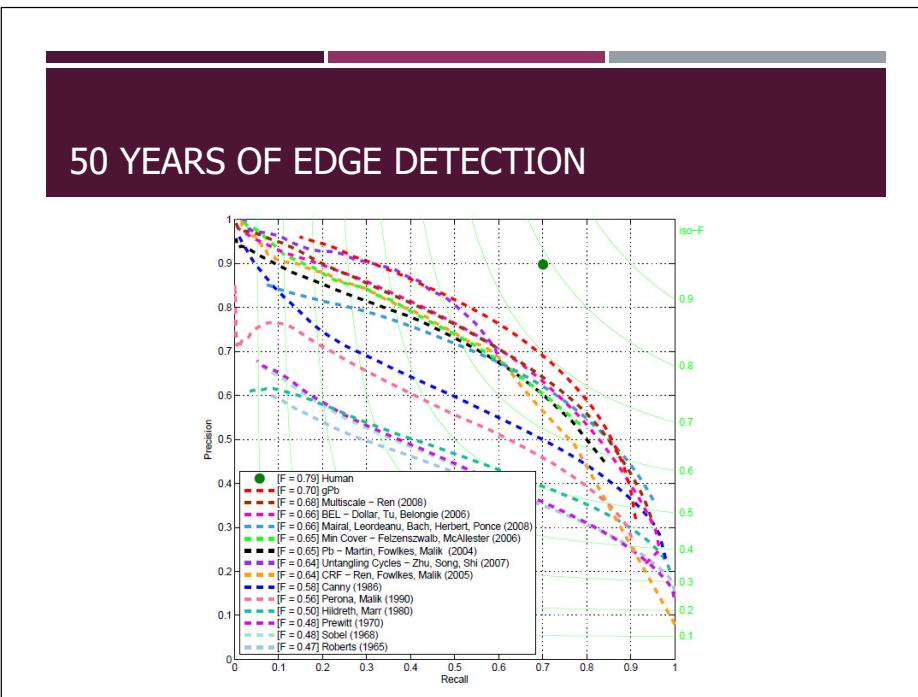


gradient magnitude



- Berkeley segmentation database:
<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/seabench/>





Edge Detection

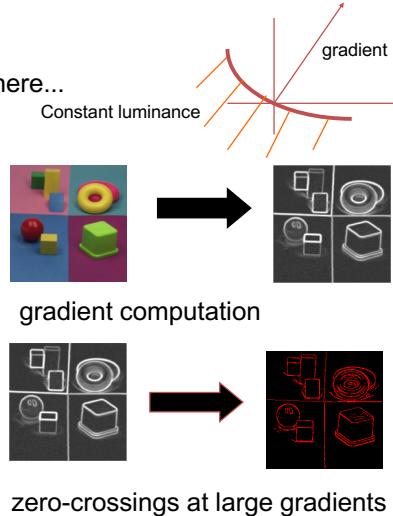
An intensity edge is defined as a point where...

the gradient is large:

$$|\nabla f(x, y)| = \sqrt{f_x^2 + f_y^2}$$

Localization - Laplacian - zero-crossing:

$$\nabla^2 f(x, y) = f_{xx} + f_{yy}$$



01-03-18

Computer Vision 1

49

Color Edge Detection

Summing the gradient magnitudes separately:

$$|\nabla C(x, y)| = \sqrt{(R_x^2 + R_y^2)} + \sqrt{(G_x^2 + G_y^2)} + \sqrt{(B_x^2 + B_y^2)}$$

or using the Euclidean metric:

$$|\nabla C(x, y)| = \sqrt{R_x^2 + R_y^2 + G_x^2 + G_y^2 + B_x^2 + B_y^2}$$

or using eigen-values: [diZzenzo86], [Sapiro96]

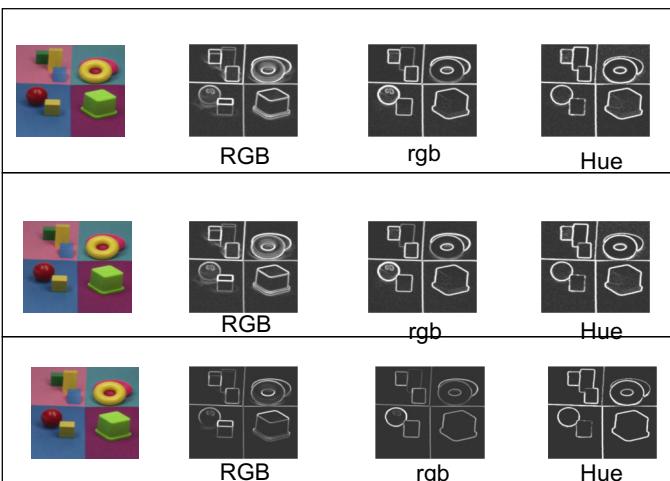
01-03-18

Computer Vision 1

50

Edge Classification

summing
gradient
magnitudes
separately:



Euclidean:

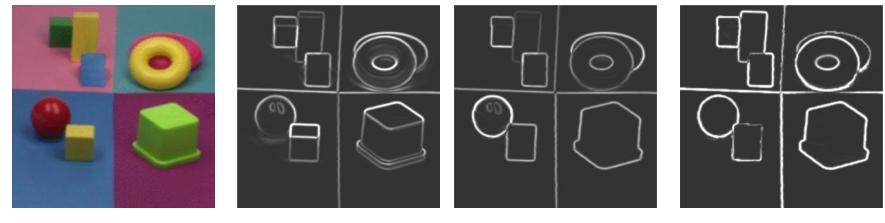
eigen-values:

01-03-18

Computer Vision 1

51

Edge Classification



if ($|\nabla C_{rgb}| \geq t_{rgb}$ & $|\nabla C_H| < t_H$) then classify as highlight edge

else

if ($|\nabla C_H| \geq t_H$) then classify as color edge

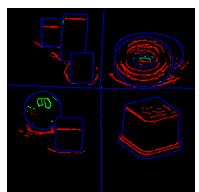
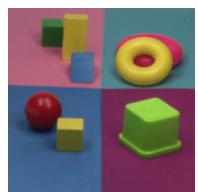
else classify as shadow/geometry edge

01-03-18

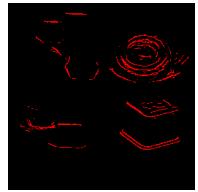
Computer Vision 1

52

Edge Classification



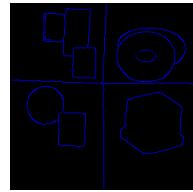
colour edge maxima by type



shadows and geometry



highlights



colour edges

01-03-18

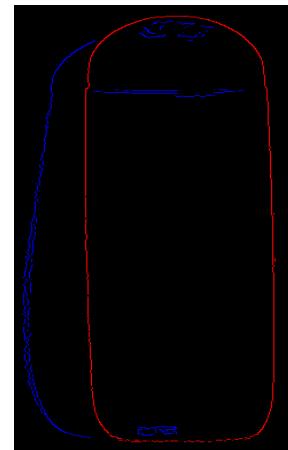
Computer Vision 1

53

Edge Classification (1)



material
shadow or geometry



01-03-18

Computer Vision 1

54

Edge Classification (2)

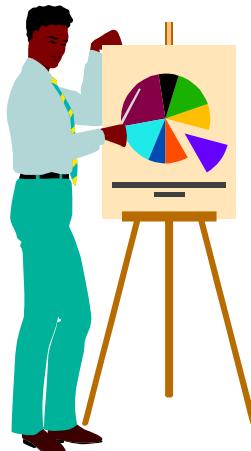


Reflectance-based Classification of Color Edges , Th. Gevers, ICCV 2003

INTEREST POINT DETECTORS

1. Edges
2. Corners
3. Blobs
4. Templates

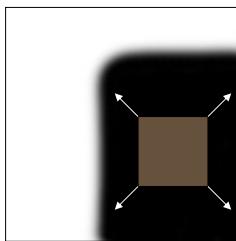
Harris Corner Detector



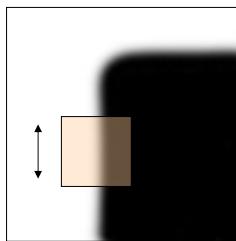
Corners as distinctive interest points

We should easily recognize the point by looking through a small window

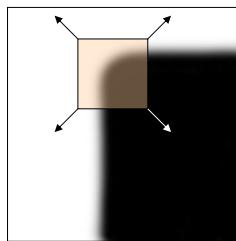
Shifting a window in *any direction* should give a *large change* in intensity



“flat” region:
no change in
all directions

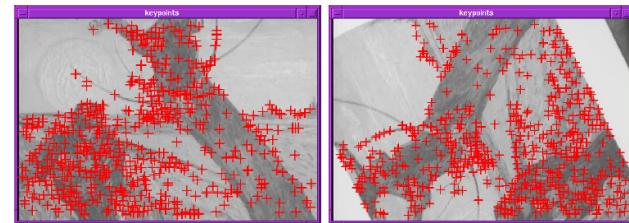


“edge”:
no change
along the edge
direction



“corner”:
significant
change in all
directions

Finding Corners



Key property: in the region around a corner, image gradient has two or more dominant directions

Corners are repeatable and **distinctive**

C.Harris and M.Stephens. ["A Combined Corner and Edge Detector."](#)
Proceedings of the 4th Alvey Vision Conference: pages 147–151.

Source: Lana Lazebnik

Harris Detector formulation

Change of intensity for the shift $[u, v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Annotations below the equation:

- Window function (points to the $w(x, y)$ term)
- Shifted intensity (points to the $I(x+u, y+v)$ term)
- Intensity (points to the $I(x, y)$ term)

Window function $w(x, y) =$

1 in window, 0 outside or Gaussian

Source: A. Efros

Source: R. Szeliski

Measuring the “Properties” of E()

$$E(u, v) = \sum [I(x-u, y-v) - I(x, y)]^2 =$$

$$= \sum_{x,y} [I(x, y) - I_x u - I_y v + \varepsilon - I(x, y)]^2 \approx$$

$$\approx \sum_{x,y} [-I_x u - I_y v]^2 =$$

$$= \sum_{x,y} I_x^2 u^2 + I_x u I_y v + I_y v I_x u + I_y^2 v^2 =$$

$$= \sum_{x,y} [u \quad v] \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix} [u \quad v]$$

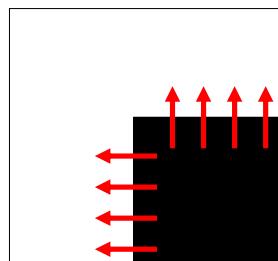
$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \sum_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix}$$

M depends on image properties

What does this matrix reveal?

First, consider an axis-aligned corner:



Harris Detector formulation

This measure of change can be approximated by:

$$E(u, v) \approx [u \quad v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Sum over image region – area we are checking for corner

$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \quad I_y]$$

Gradient with respect to x, times gradient with respect to y

What does this matrix reveal?

First, consider an axis-aligned corner:

$$M = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

This means dominant gradient directions align with x or y axis

If either λ is close to 0, then this is **not** a corner, so look for locations where both are large.

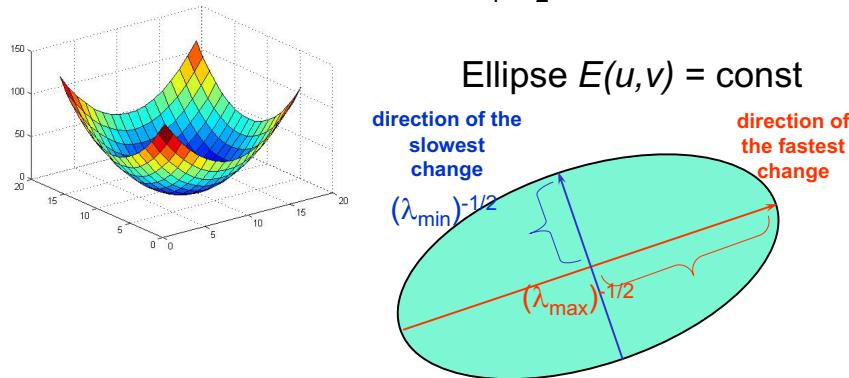
What if we have a corner that is not aligned with the image axes?

Quadratic Form and its Eigenvalue

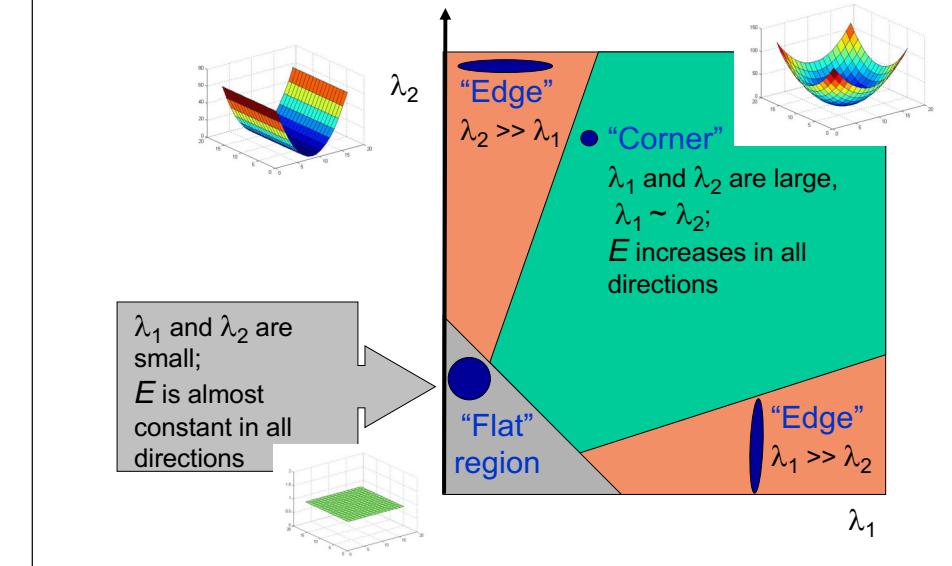
$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = U^T \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} U =$$

λ_1, λ_2 – eigenvalues of M



Classification of Image Points using Eigenvalues of M :



Harris Detector: Mathematics

Measure of corner response:

$$R = \det M - k (\text{trace } M)^2$$

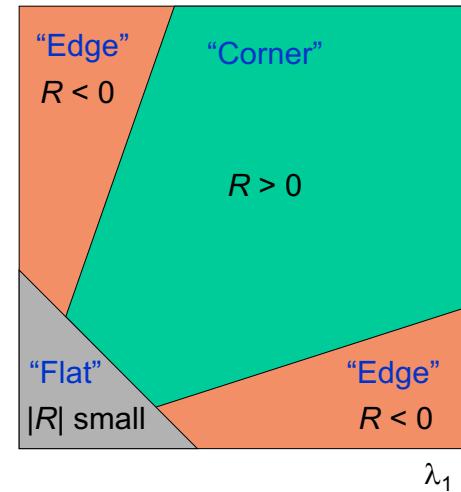
$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

(k – empirical constant, $k = 0.04-0.06$)

Harris Detector

- R depends only on eigenvalues of M
- R is large for a corner
- R is negative with large magnitude for an edge
- $|R|$ is small for a flat region



$$R(x_0, y_0) = \det M - k (\text{trace } M)^2$$

Summary of the Harris detector

1. Compute x and y derivatives of image

$$I_x = G_\sigma^x * I \quad I_y = G_\sigma^y * I$$

2. Compute products of derivatives at every pixel

$$I_{x2} = I_x \cdot I_x \quad I_{y2} = I_y \cdot I_y \quad I_{xy} = I_x \cdot I_y$$

3. Compute the sums of the products of derivatives at each pixel

$$S_{x2} = G_{\sigma_I} * I_{x2} \quad S_{y2} = G_{\sigma_I} * I_{y2} \quad S_{xy} = G_{\sigma_I} * I_{xy}$$

4. Define at each pixel (x, y) the matrix

$$H(x, y) = \begin{bmatrix} S_{x2}(x, y) & S_{xy}(x, y) \\ S_{xy}(x, y) & S_{y2}(x, y) \end{bmatrix}$$

5. Compute the response of the detector at each pixel

$$R = \text{Det}(H) - k(\text{Trace}(H))^2$$

6. Threshold on value of R. Compute nonmax suppression.

01-03-18

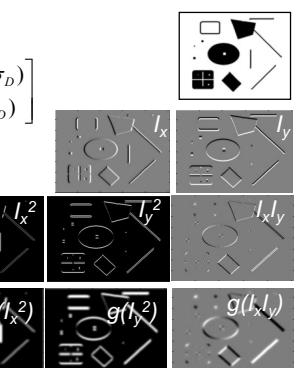
Computer Vision 1

69

Harris Detector [Harris88]

$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

1. Image derivatives

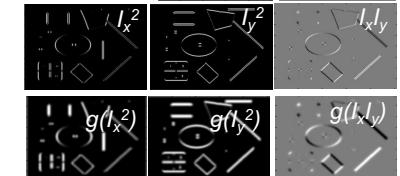


2. Square of derivatives

$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

3. Gaussian filter $g(s_I)$

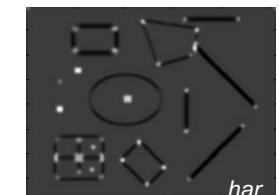


4. Cornerness function – both eigenvalues are strong

$$\begin{aligned} \text{har} &= \det[\mu(\sigma_I, \sigma_D)] - \alpha[\text{trace}(\mu(\sigma_I, \sigma_D))^2] = \\ &= g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - \alpha[g(I_x^2) + g(I_y^2)]^2 \end{aligned}$$

5. Non-maxima suppression

Computer Vision 1
Slide: Derek Hoiem

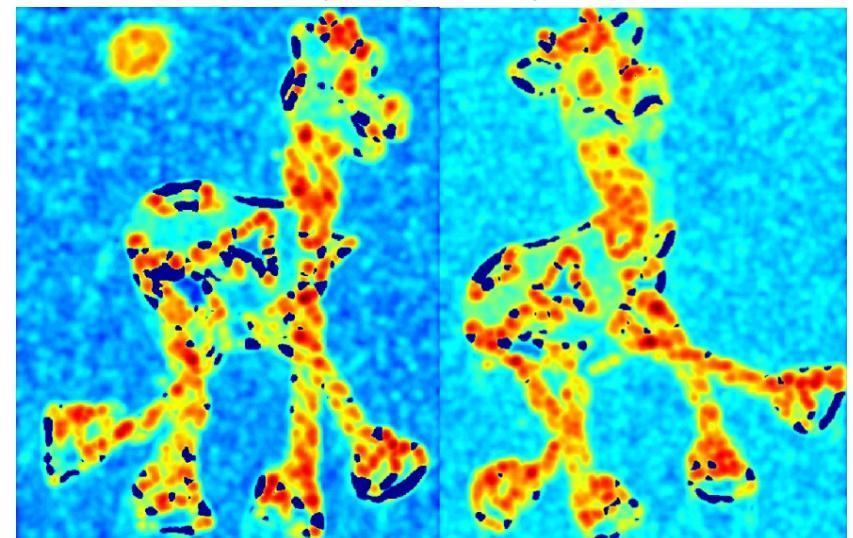


Harris Detector: Workflow



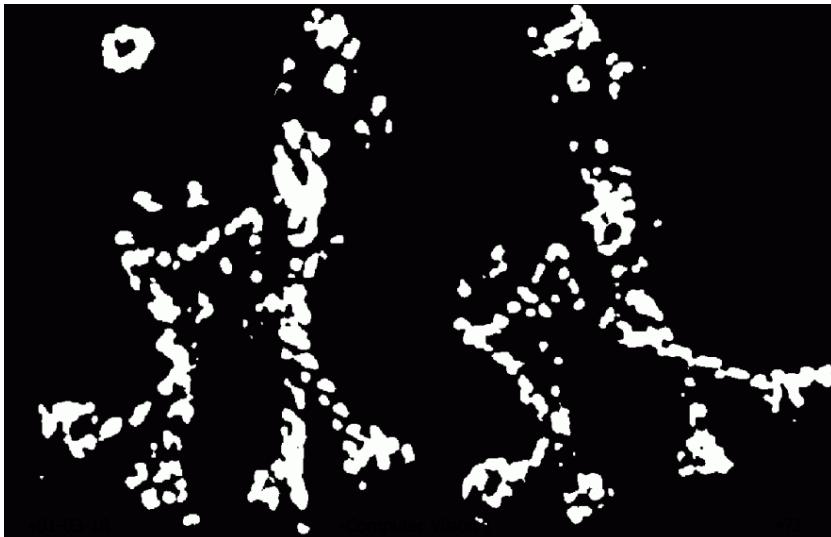
Harris Detector: Workflow

Compute corner response R



Harris Detector: Workflow

Find points with large corner response: $R > \text{threshold}$



Harris Detector: Workflow

Take only the points of local maxima of R

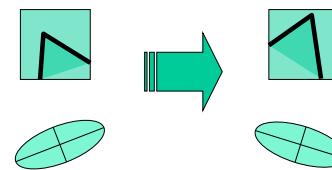


Harris Detector: Workflow



Harris Detector: Properties

Rotation invariance



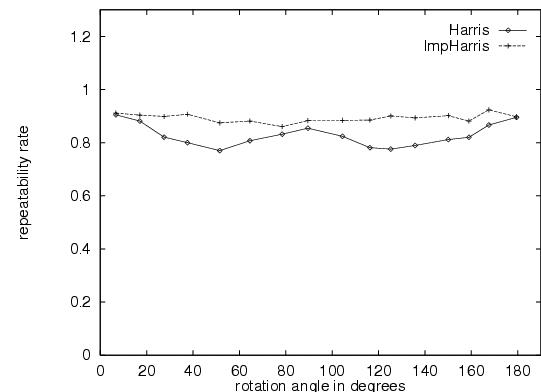
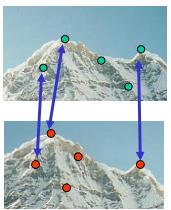
Ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner response R is invariant to image rotation

Harris Detector: Properties

Repeatability rate:

$$\frac{\# \text{ correspondences}}{\# \text{ possible correspondences}}$$



01-03-18

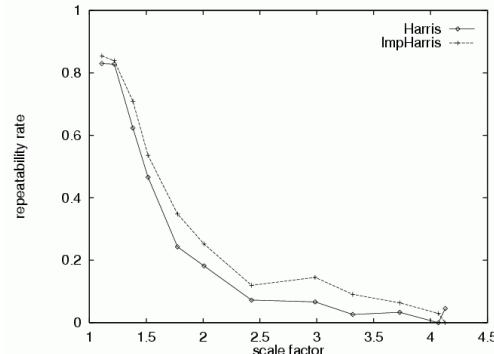
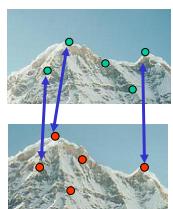
Computer Vision 1

77

Harris Detector: Properties

Repeatability rate:

$$\frac{\# \text{ correspondences}}{\# \text{ possible correspondences}}$$



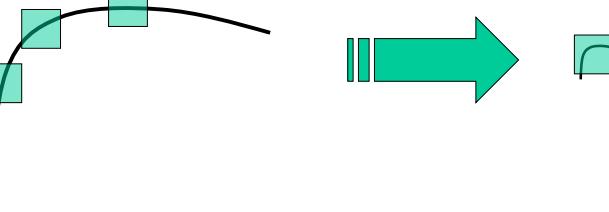
01-03-18

Computer Vision 1

79

Harris Detector: Properties

Not invariant to image scale



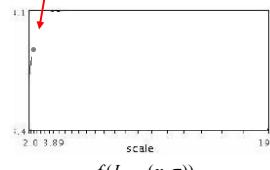
All points will be
classified as edges

Corner !

How can we detect **scale invariant**
interest points?

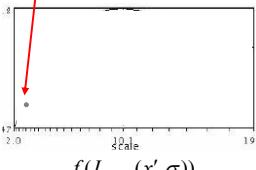
Automatic Scale Selection

Function responses for increasing scale (scale signature)



$f(I_{i_1...i_m}(x, \sigma))$

K. Grauman, B. Leibe



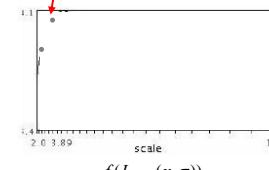
$f(I_{i_1...i_m}(x', \sigma))$

K. Grauman, B. Leibe

81

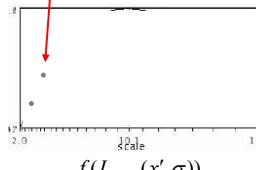
Automatic Scale Selection

Function responses for increasing scale (scale signature)



$f(I_{i_1...i_m}(x, \sigma))$

K. Grauman, B. Leibe



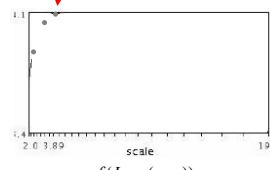
$f(I_{i_1...i_m}(x', \sigma))$

K. Grauman, B. Leibe

82

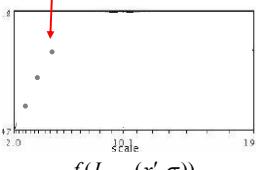
Automatic Scale Selection

Function responses for increasing scale (scale signature)



$f(I_{i_1...i_m}(x, \sigma))$

K. Grauman, B. Leibe



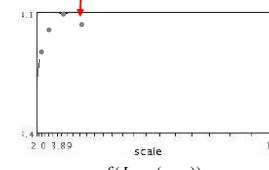
$f(I_{i_1...i_m}(x', \sigma))$

K. Grauman, B. Leibe

83

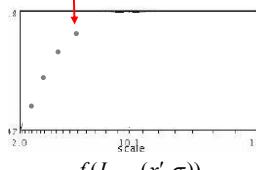
Automatic Scale Selection

Function responses for increasing scale (scale signature)



$f(I_{i_1...i_m}(x, \sigma))$

K. Grauman, B. Leibe



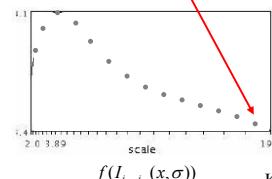
$f(I_{i_1...i_m}(x', \sigma))$

K. Grauman, B. Leibe

84

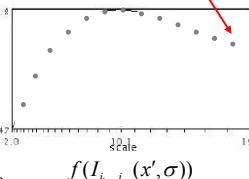
Automatic Scale Selection

Function responses for increasing scale (scale signature)



$f(I_{i_1...i_m}(x, \sigma))$

K. Grauman, B. Leibe
K. Grauman, B. Leibe

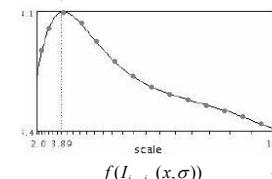


$f(I_{i_1...i_m}(x', \sigma))$

85

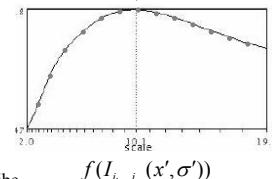
Automatic Scale Selection

Function responses for increasing scale (scale signature)



$f(I_{i_1...i_m}(x, \sigma))$

K. Grauman, B. Leibe
K. Grauman, B. Leibe



$f(I_{i_1...i_m}(x', \sigma'))$

86

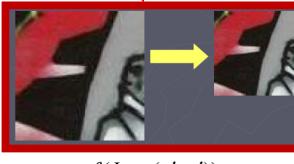
Automatic Scale Selection

Function responses for increasing scale (scale signature)



$f(I_{i_1...i_m}(x, \sigma))$

K. Grauman, B. Leibe
K. Grauman, B. Leibe



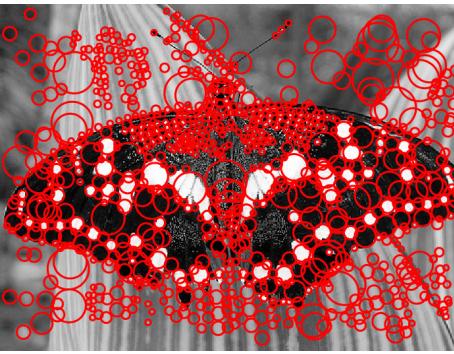
$f(I_{i_1...i_m}(x', \sigma'))$

87

INTEREST POINT DETECTORS

1. Edges
2. Corners
3. Blobs
4. Templates

Scale-space Blob Detector: Example



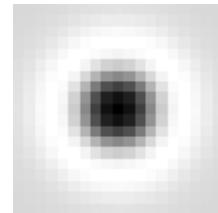
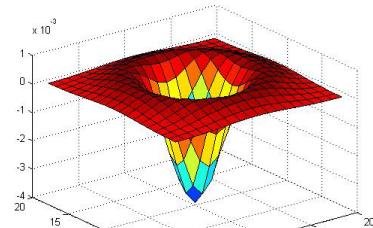
01-03-18

Computer Vision 1

89
Image credit: Lana Lazebnik

Laplacian of Gaussian

- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D



$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

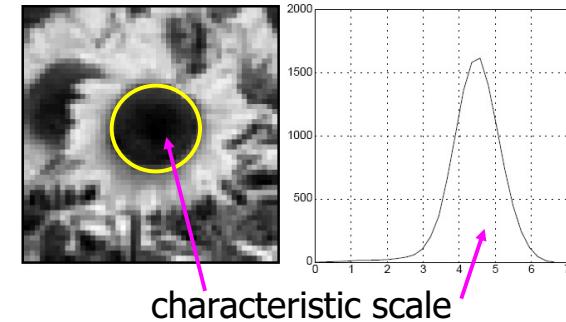
01-03-18

Computer Vision 1

91

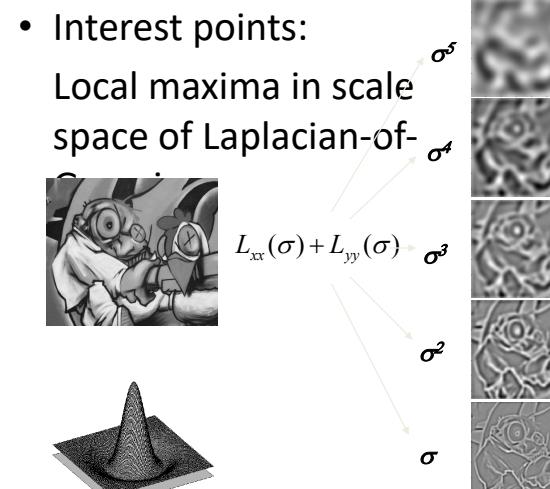
Characteristic scale

- We define the *characteristic scale* as the scale that produces peak of Laplacian response



T. Lindeberg (1998). ["Feature detection with automatic scale selection."](#)
International Journal of Computer Vision **30** (2): pp 77--116. Source: Lana Lazebnik

Laplacian-of-Gaussian (LoG)



K. Grauman, B. Leibe
K. Grauman, B. Leibe

Scale-space blob detector: Example



Source: Lana Lazebnik

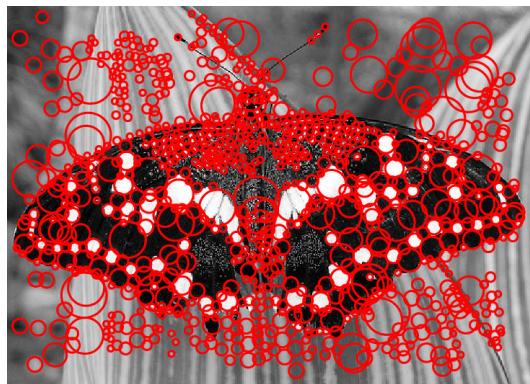
Scale-space blob detector: Example



$\sigma = 11.9912$

Source: Lana Lazebnik

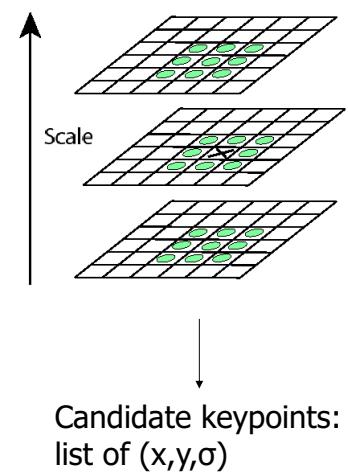
Scale-space blob detector: Example



Source: Lana Lazebnik

Key point localization with DoG

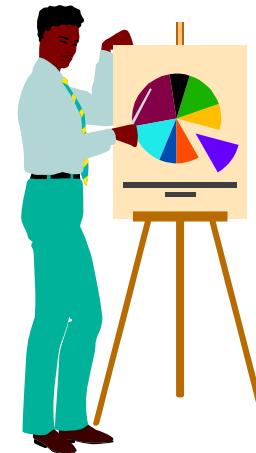
- Detect maxima of difference-of-Gaussian (DoG) in scale space
- Then reject points with low contrast (threshold)
- Eliminate edge responses



INTEREST POINT DETECTORS

1. Edges
2. Corners
3. Blobs
4. Templates

Template Matching

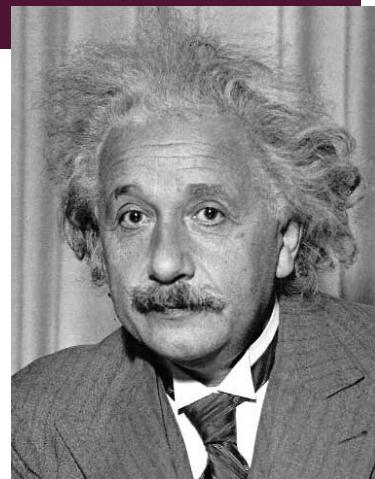


TEMPLATE MATCHING

- Goal: find in image

Main challenge:

- What is a good similarity or distance measure between two patches?
 - Sum Square Difference
 - Normalized Cross Correlation
 - (Matlab: normxcorr2)



Template Matching

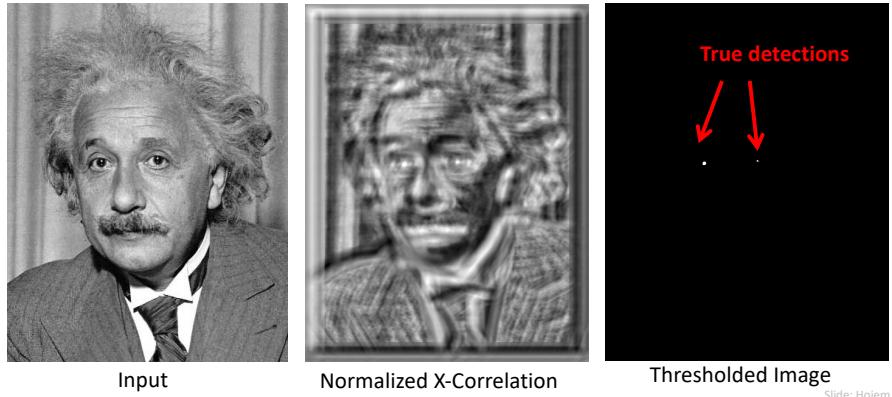
- Goal: find in image
- SSD $h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$



Slide: Holm

Template Matching

- Goal: find  in image
- Method 2: Normalized cross-correlation



WHERE WILL INTEREST POINTS COME FROM?

1. Edges
2. Corners
3. Blobs
4. Templates

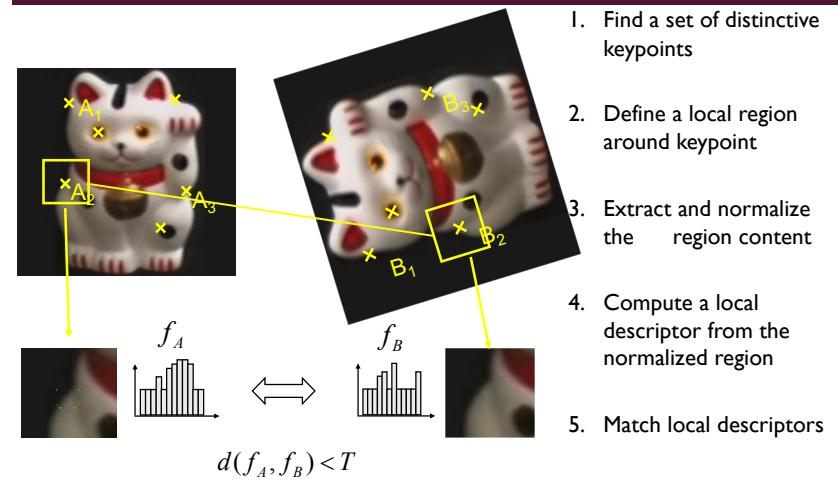
Comparison of Keypoint Detectors

Table 7.1 Overview of feature detectors.

Feature Detector	Corner	Blob	Region	Rotation invariant	Scale invariant	Affine invariant	Repeatability	Localization accuracy	Robustness	Efficiency
Harris	✓			✓			+++	+++	+++	++
Hessian		✓		✓			++	++	++	+
SUSAN	✓			✓			++	++	++	+++
Harris-Laplace	✓	(✓)		✓	✓		+++	+++	++	+
Hessian-Laplace	(✓)	✓		✓	✓		+++	+++	+++	+
DoG	(✓)	✓		✓	✓		++	++	++	++
SURF	(✓)	✓		✓	✓		++	++	++	+++
Harris-Affine	✓	(✓)		✓	✓	✓	+++	+++	++	++
Hessian-Affine	(✓)	✓		✓	✓	✓	+++	+++	+++	++
Salient Regions	(✓)	✓		✓	✓	(✓)	+	+	++	+
Edge-based	✓			✓	✓	✓	+++	+++	+	+
MSER			✓	✓	✓	✓	+++	+++	++	+++
Intensity-based			✓	✓	✓	✓	++	++	++	++
Superpixels	✓		✓	(✓)	(✓)	(✓)	+	+	+	+

From a study of Tuytelaars & Mikolajczyk

OVERVIEW OF KEYPOINT MATCHING

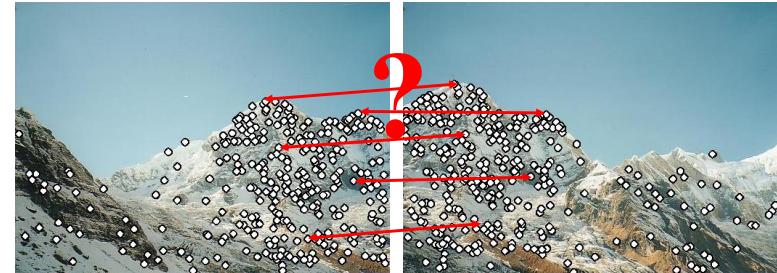


How to describe interest points?

Local descriptors

- We know how to detect points
- Next question:

How to describe them for matching?



Point descriptor should be:

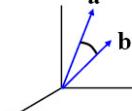
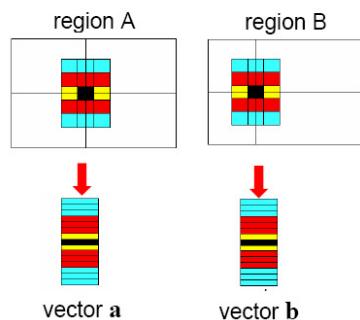
1. Invariant
2. Distinctive

Local descriptors

- Simplest descriptor: list of intensities within a patch.
- What is this going to be invariant to?

Write regions as vectors

$$A \rightarrow a, B \rightarrow b$$



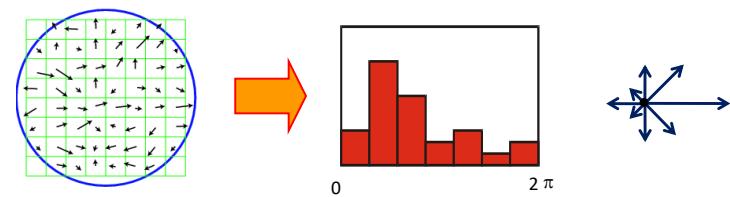
Feature descriptors

Disadvantage of patches as descriptors:

- Small shifts can affect matching score a lot



Solution: histograms



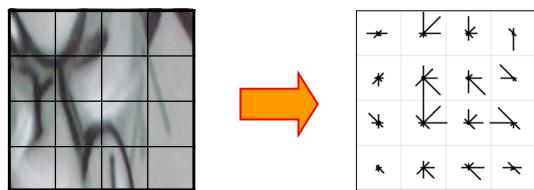
Source: Lana Lazebnik

Feature descriptors: SIFT

Scale Invariant Feature Transform

Descriptor computation:

- Divide patch into 4×4 sub-patches: 16 cells
- Compute histogram of gradient orientations (8 reference angles) for all pixels inside each sub-patch
- Resulting descriptor: $4 \times 4 \times 8 = 128$ dimensions



David G. Lowe. ["Distinctive image features from scale-invariant keypoints."](#) IJCV 60 (2), pp. 91-110, 2004.

Source: Lana Lazebnik

Rotation Invariant Descriptors

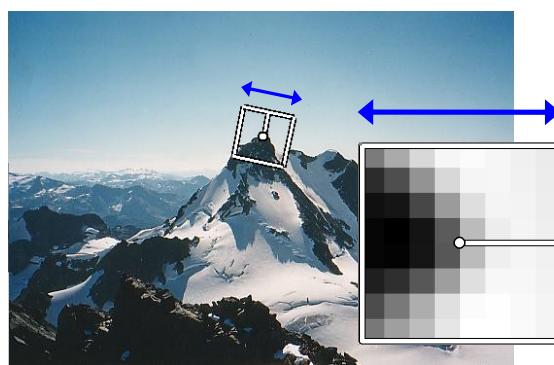


Image from Matthew Brown

Rotation Invariant Descriptors

- Find local orientation

Dominant direction of gradient for the image patch



- Rotate patch according to this angle

This puts the patches into a canonical orientation.

Feature descriptors: SIFT

Extraordinarily robust matching technique

- Can handle changes in viewpoint
 - Up to about 60 degree out of plane rotation
- Can handle significant changes in illumination
 - Sometimes even day vs. night (below)
- Fast and efficient—can run in real time
- Lots of code available
 - http://people.csail.mit.edu/albert/ladypack/wiki/index.php/Known_implementations_of_SIFT



Steve Seitz

Working with SIFT descriptors

- One image yields:
 - n 128-dimensional descriptors: each one is a histogram of the gradient orientations within a patch
 - [n x 128 matrix]
 - n scale parameters specifying the size of each patch
 - [n x 1 vector]
 - n orientation parameters specifying the angle of the patch
 - [n x 1 vector]
 - n 2d points giving positions of the patches
 - [n x 2 matrix]

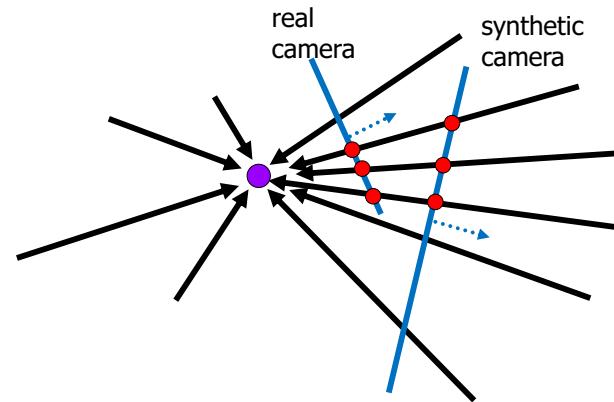


Building a Panorama



How to match interest points?

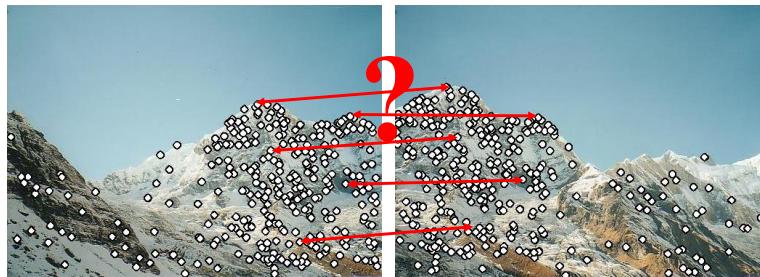
Panoramas: generating synthetic views



Can generate any synthetic camera view
as long as it has the **same center of projection!**

FEATURE MATCHING

We know how to detect and describe good points

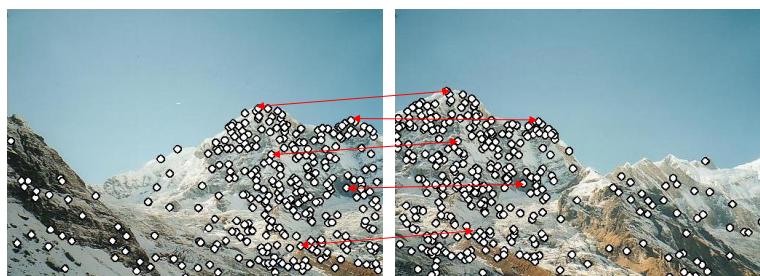


Given a feature in Img1, how to find the best match in Img2?

- Define distance function that compares two descriptors
- Test all the features in I2, find the one with min distance

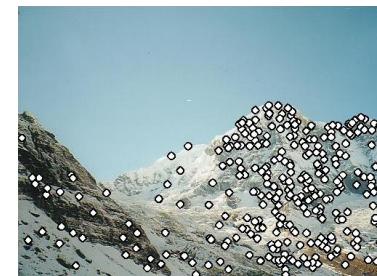
Building a Panorama

1. Detect feature points in both images
2. Find corresponding pairs



Building a Panorama

- 1) Detect feature points in both images



Building a Panorama

1. Detect feature points in both images
2. Find corresponding pairs
3. Find a parametric transformation (e.g. homography)
4. Warp (right image to left image)



$$\forall \text{ matching pair } \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}_{\text{left}} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}_{\text{right}}$$

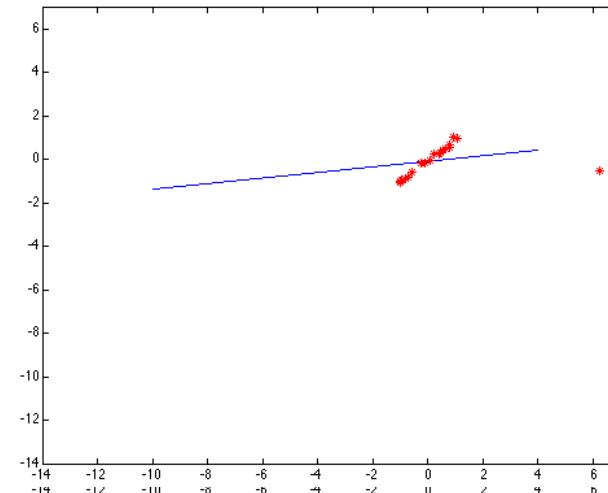
Outliers

- **Outliers** can hurt the quality of our parameter estimates, e.g.,
 - an erroneous pair of matching points from two images
 - an edge point that is noise, or doesn't belong to the line we are fitting.



Grauman

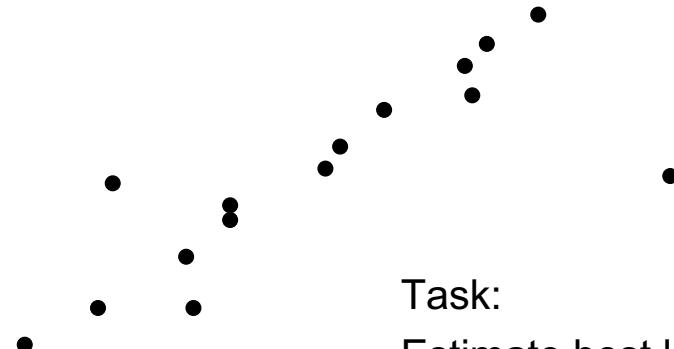
Example: Outliers in Linear Regression



RANSAC

- RANdom Sample Consensus
- Approach: we want to avoid the impact of outliers, so let's look for “inliers”, and use those only.
- Intuition: if an outlier is chosen to compute the current fit, then the resulting line won't have much support from rest of the points.

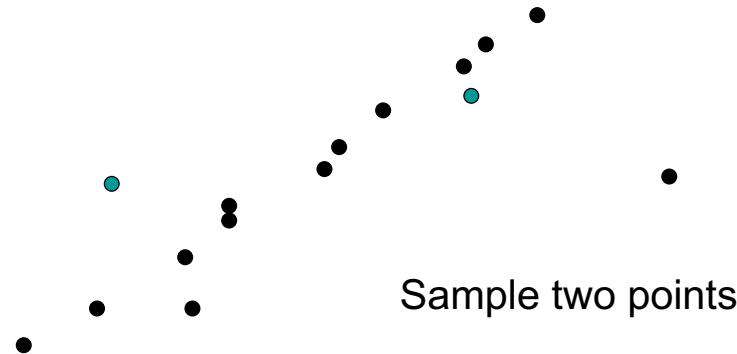
RANSAC Line Fitting Example



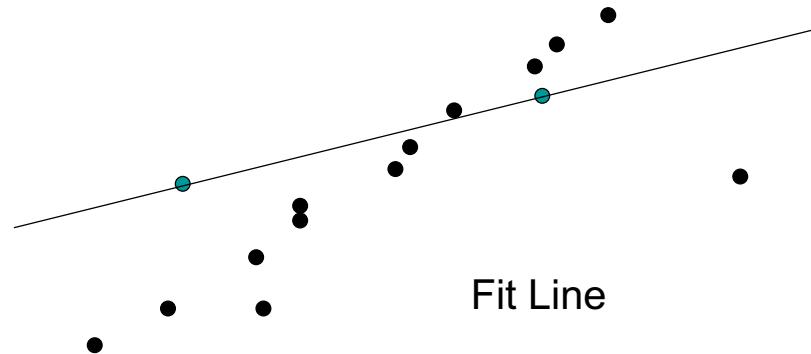
Task:
Estimate best line

Slide credit: Jinxiang Chai, CMU

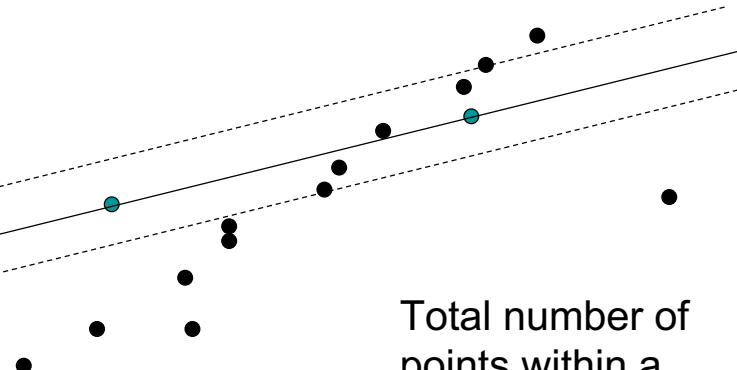
RANSAC Line Fitting Example



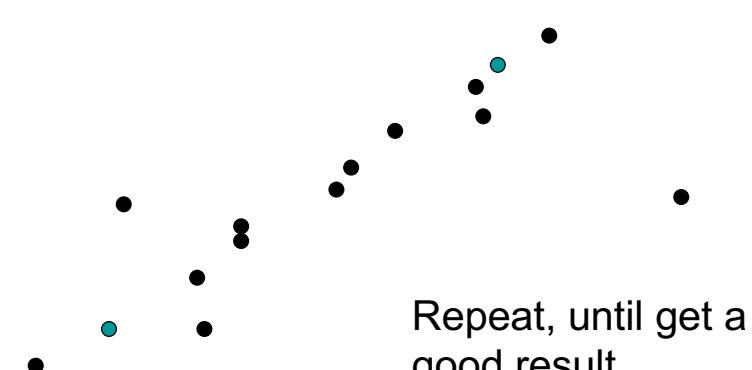
RANSAC Line Fitting Example



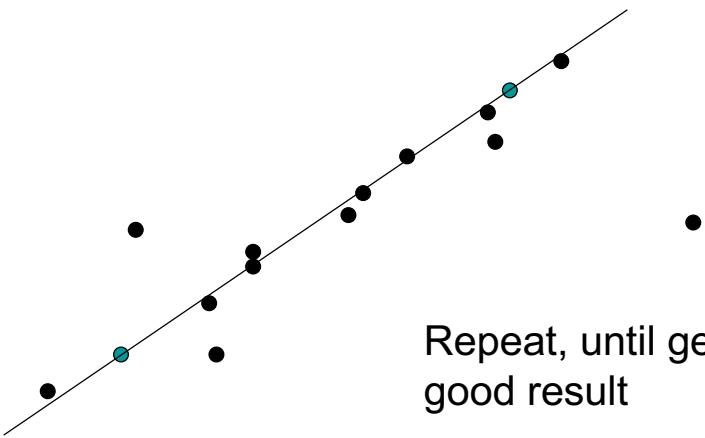
RANSAC Line Fitting Example



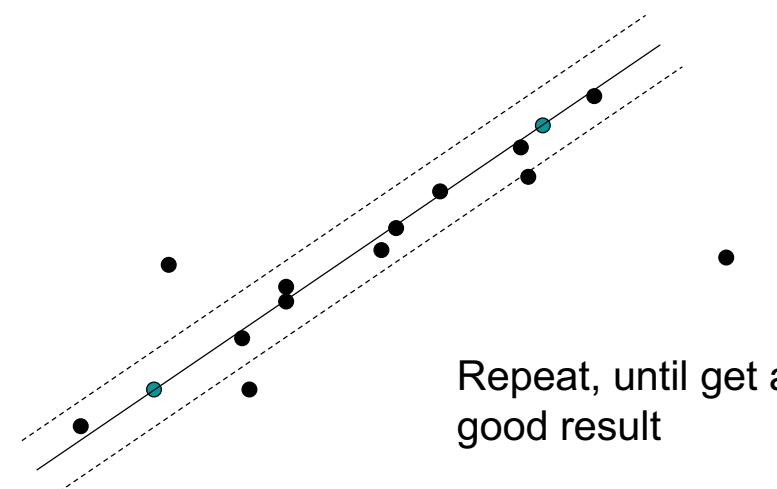
RANSAC Line Fitting Example



RANSAC Line Fitting Example



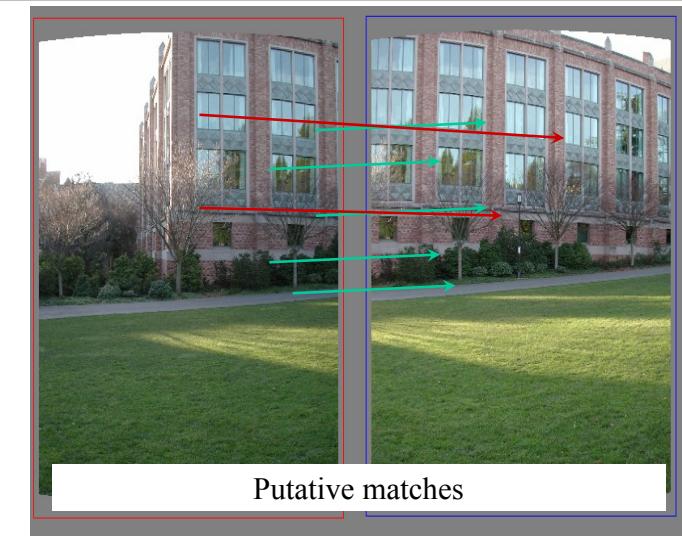
RANSAC Line Fitting Example



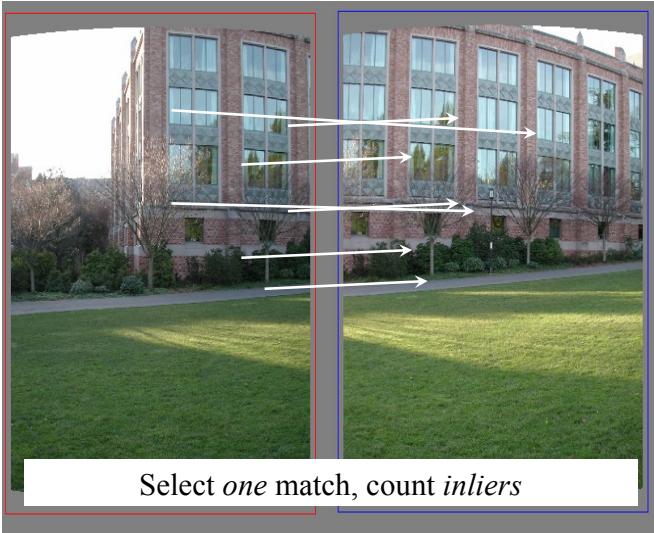
RANSAC for Estimating Homography

- RANSAC loop:
 1. Select four feature pairs (at random)
 2. Compute homography H (exact)
 3. Compute *inliers* where $SSD(p'_i, H p_i) < \varepsilon$
 4. Keep largest set of inliers
 5. Re-compute least-squares H estimate on all of the inliers

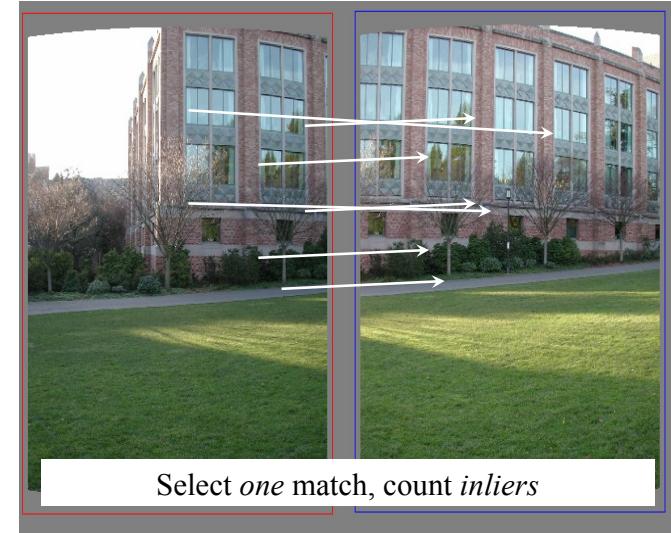
RANSAC example: Translation



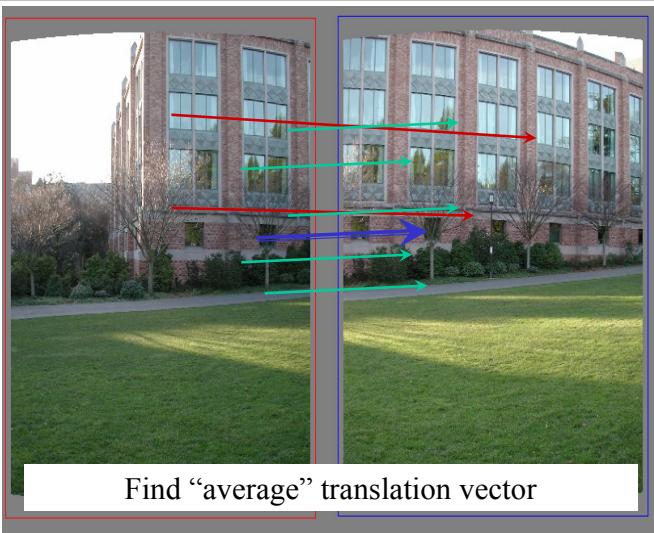
RANSAC example: Translation

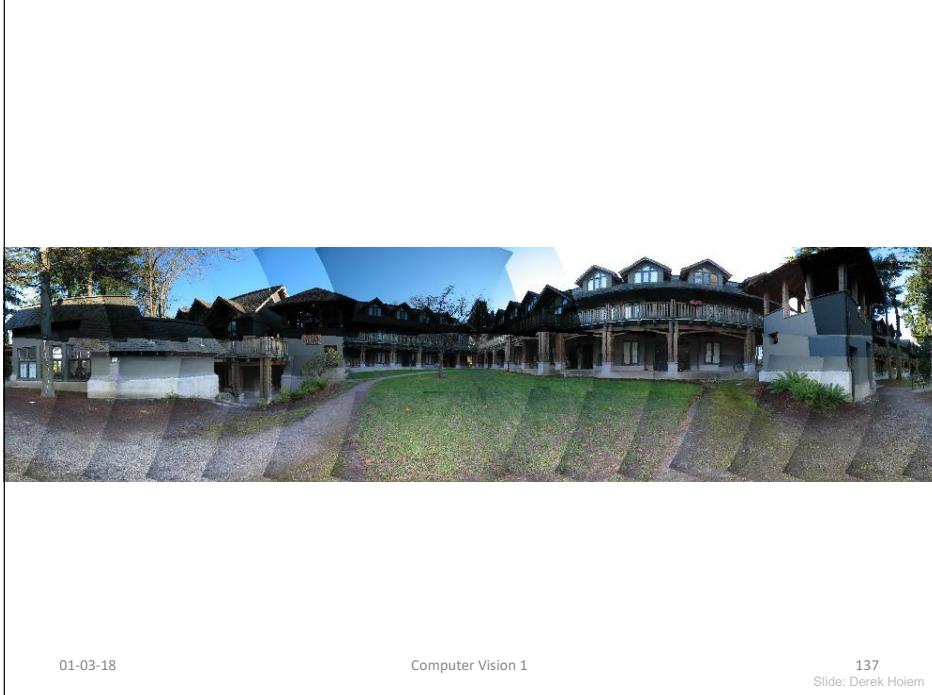


RANSAC example: Translation



RANSAC example: Translation

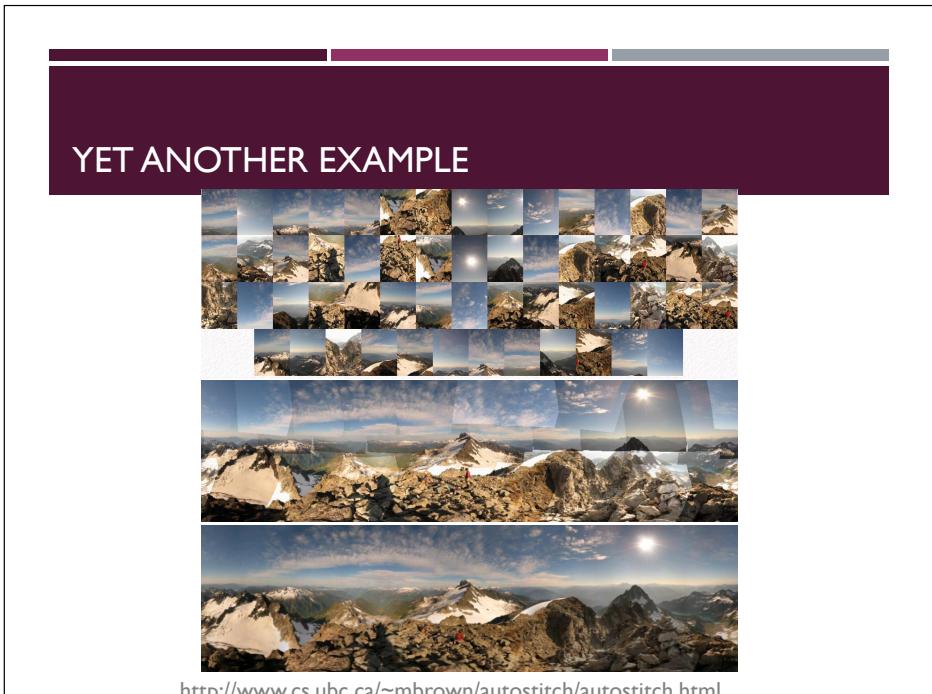




01-03-18

Computer Vision 1

137
Slide: Derek Hoiem



<http://www.cs.ubc.ca/~mbrown/autostitch/autostitch.html>

Multi-band Blending

- Burt & Adelson 1983
 - Blend frequency bands over range $\propto I$



01-03-18

Computer Vision 1

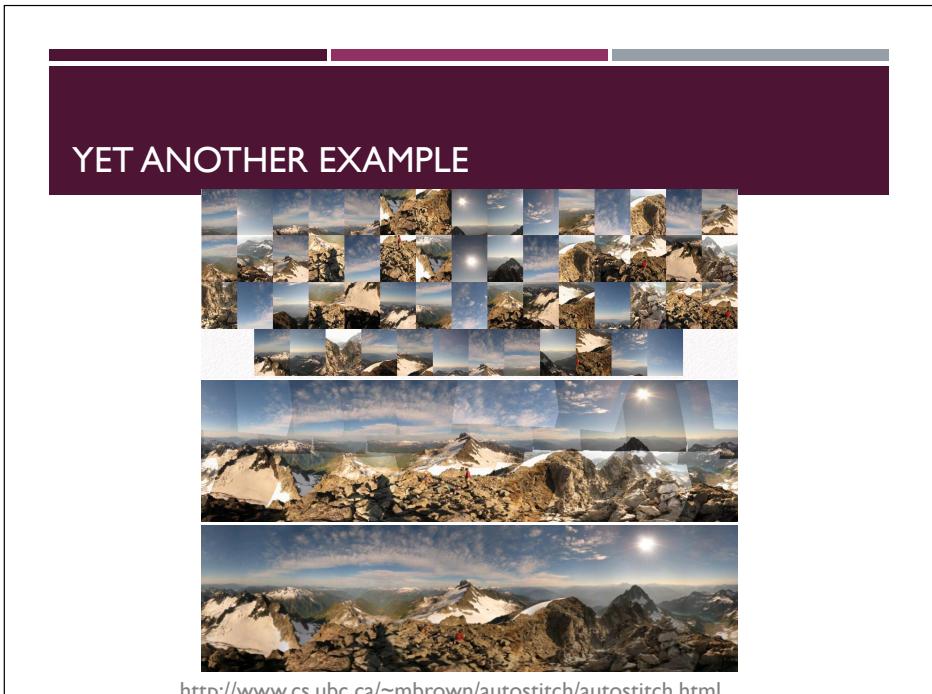
137
Slide: Derek Hoiem



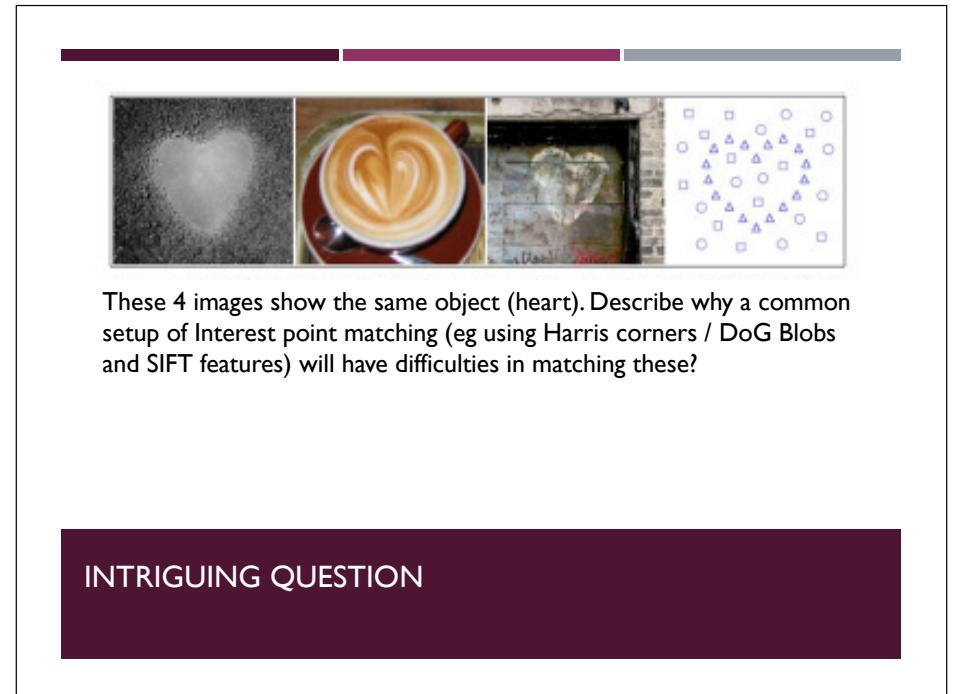
01-03-18

Computer Vision 1

138
Slide: Derek Hoiem



<http://www.cs.ubc.ca/~mbrown/autostitch/autostitch.html>



These 4 images show the same object (heart). Describe why a common setup of Interest point matching (eg using Harris corners / DoG Blobs and SIFT features) will have difficulties in matching these?

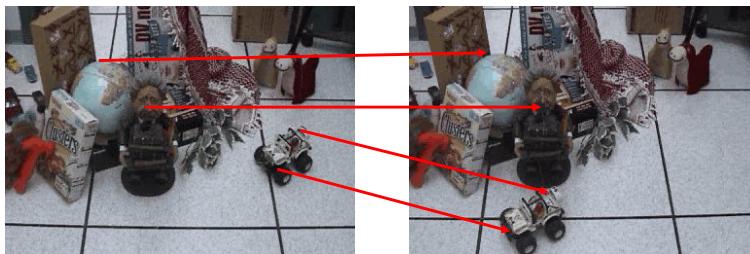
INTRIGUING QUESTION

TODAY'S CLASS

- Image Similarity
- Feature Detection
- Motion

OPTICAL FLOW

Where did each pixel in image 1 go to in image 2



MOTION & VISUAL TRACKING

- Optical Flow
- Template tracking
- Mean-shift tracking
- Foreground-Background Separation Tracking
- Tracking by Detection

OPTICAL FLOW (EXAMPLE VIDEO)



MOTIVATION

Stabilization

- Visual Motion can be annoying
 - Camera instabilities, jitter
 - Camera movement
- Measure it; remove it (stabilize)

Scene Recognition

- Visual Motion reveals spatial layout
 - Motion parallax

Object recognition

- Visual Motion indicates dynamics
 - How many moving objects are there.
 - Which directions are they moving in.
 - How fast are they moving.
- Track objects and analyze trajectories

Action Recognition

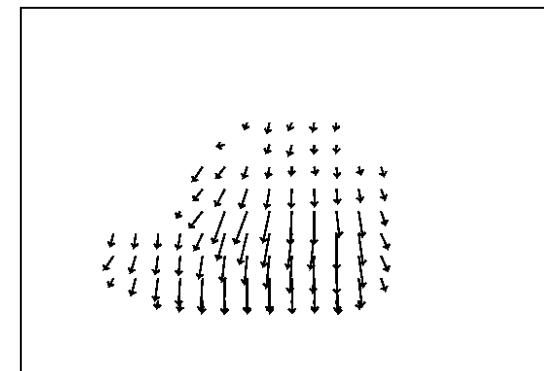
- Can we recognize their type of motion (e.g. walking, running, etc.).

Computer Vision 1

01-03-18

145

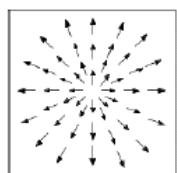
FLOW FIELD



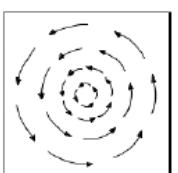
Pierre Kornprobst's Demo
01-03-18

146

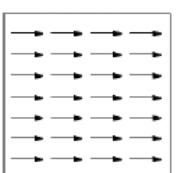
Examples of Motion fields



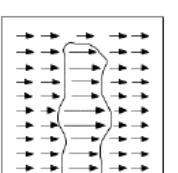
Forward motion



Rotation



Horizontal translation



Closer objects appear to move faster!!

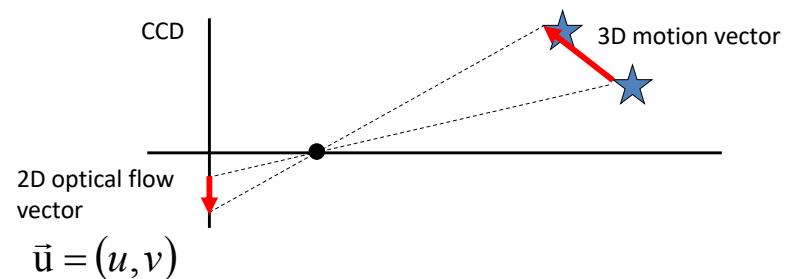
01-03-18

Computer Vision 1

147

Motion Field & Optical Flow Field

- Motion Field = Real world 3D motion
- Optical Flow Field = Projection of the motion field onto the 2d image



01-03-18

Computer Vision 1

148

The Optical Flow Field

- Goal:
Find for each pixel a velocity vector which says:
 - In which direction it is moving
 - How quickly is the pixel moving across the image

01-03-18

Computer Vision 1

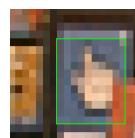
149

Image motion

How do we determine correspondences?



I



J

Assume all change between frames is due to **motion!**

Patch-based motion estimation

Estimating Optical Flow

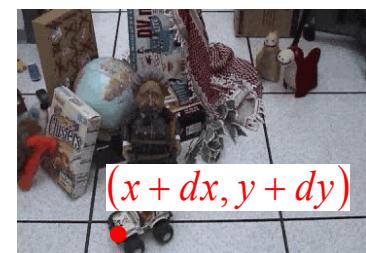
Assume the image intensity I is constant

Time = t



(x, y)

Time = $t+dt$



$(x + dx, y + dy)$

$$I(x, y, t) = I(x + dx, y + dy, t + dt)$$

01-03-18

Computer Vision 1

152

Brightness Constancy Equation

$$I(x, y, t) = I(x + dx, y + dy, t + dt)$$

First order Taylor Expansion

$$= I(x, y, t) + \frac{\partial I}{\partial x} dx + \frac{\partial I}{\partial y} dy + \frac{\partial I}{\partial t} dt$$

Assumption?

Simplify notations:

$$I_x dx + I_y dy + I_t dt = 0$$

Divide by dt and denote:

$$u = \frac{dx}{dt} \quad v = \frac{dy}{dt}$$

$$I_x u + I_y v = -I_t$$

Problem: One equation, two unknowns

01-03-18

Computer Vision 1

153

Conditions for Solvability

- Optimal (u, v) satisfies Lucas-Kanade equation

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$$A^T A \quad A^T b$$

- When is This Solvable?

- $A^T A$ should be invertible
- $A^T A$ should not be too small due to noise
 - eigenvalues λ_1 and λ_2 of $A^T A$ should not be too small
- $A^T A$ should be well-conditioned
 - λ_1 / λ_2 should not be too large (λ_1 = larger eigenvalue)

01-03-18

Computer Vision 1

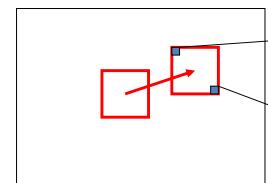
155

* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

Local Smoothness: Lucas Kanade

$$I_x u + I_y v = -I_t \rightarrow \begin{bmatrix} I_x & I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = -I_t$$

Assume constant (u, v) in small neighborhood



$$\begin{bmatrix} I_{x1} & I_{y1} \\ I_{x2} & I_{y2} \\ \vdots & \vdots \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_{t1} \\ I_{t2} \\ \vdots \end{bmatrix}$$

$$A \vec{u} = b$$

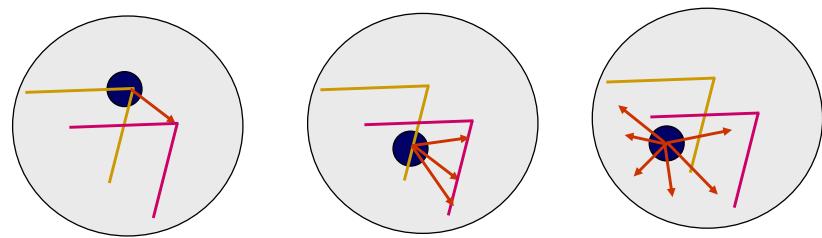
01-03-18

Computer Vision 1

154

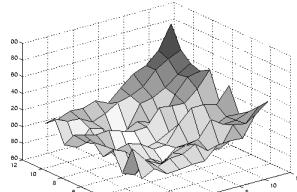
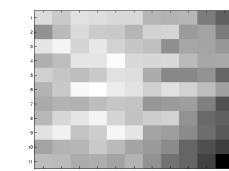
Local Patch Analysis

- How *certain* are the motion estimates?



Szelisk

Low Texture Region



$$\sum \nabla I (\nabla I)^T$$

- gradients have small magnitude
- small I_1 , small I_2

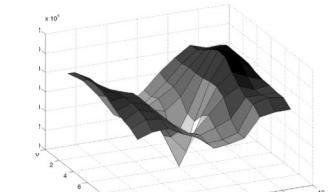
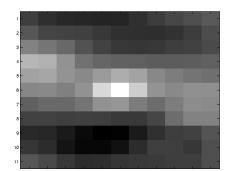
01-03-18

Computer Vision 1

157

* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

High Textured Region



$$\sum \nabla I (\nabla I)^T$$

- gradients are different, large magnitudes
- large I_1 , large I_2

01-03-18

Computer Vision 1

158

* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

Small Motion Assumption



- Is this motion small enough?
 - Probably not—it's much larger than one pixel (Taylor expansion doesn't hold)
 - How might we solve this problem?

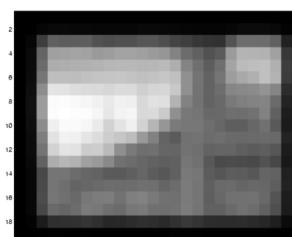
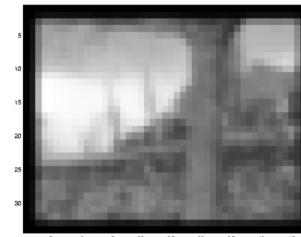
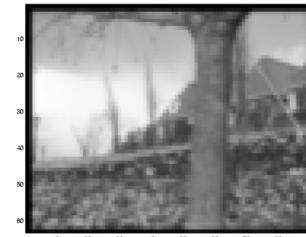
01-03-18

Computer Vision 1

159

* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

Reduce the Resolution

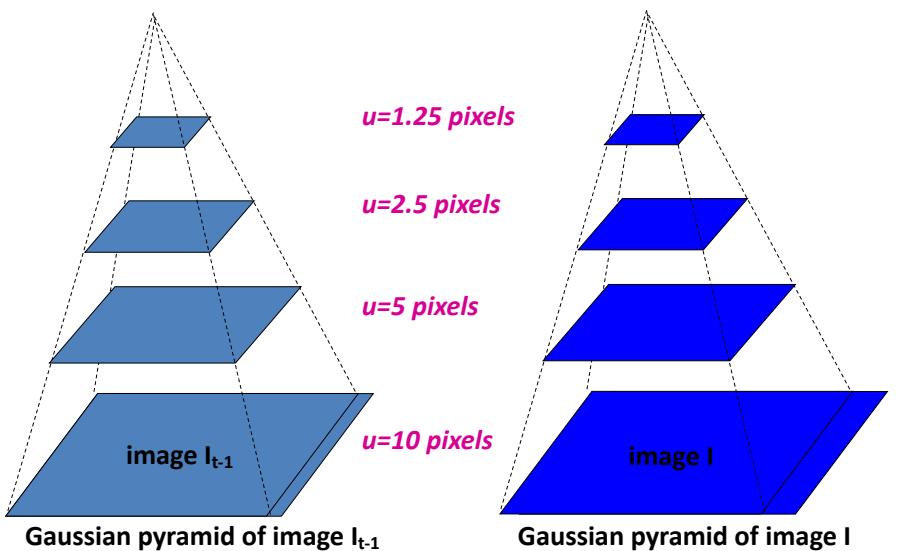


01-03-18

Computer Vision 1

160

Coarse-to-fine Optical Flow



OPTICAL FLOW (EXAMPLE VIDEO)



<https://www.youtube.com/watch?v=Wgaukp4ZJ74>

OPTICAL FLOW (RECAP)

What are the (hidden) assumption behind this flow computation?

- Finding corresponding key-points in subsequent frames
- Small motion between frames
- Intensity/Color constancy between frames

Steps

1. Find corresponding patches
2. Compute flow over small neighborhood
 - I. From coarse to fine
 - II. Using Lucas-Kanade Patch Translation

OTHER BREAK-DOWNS

Brightness constancy is **not** satisfied

→ Correlation based methods

A point does **not** move like its neighbors

what is the ideal window size?

→ Regularization based methods

MOTION & VISUAL TRACKING

- Optical Flow
- Template tracking
- Mean-shift tracking
- Foreground-Background Separation Tracking
- Tracking by Detection

Not discussed

QUESTIONS?



*The important thing is not to stop questioning;
curiosity has its own reason for existing.*

— Albert Einstein

TODAY'S CLASS

- **Image Similarity**
- **Feature Detection**
- **Motion**

- Deadline lab assignment: Gabor & Gaussian Filters
- Lab session (March 1): Optical Flow & Harris
- Lecture 4 (Tuesday): Features & Classification

SLIDE CREDITS

- Theo Gevers
- Trevor Darrell (<https://people.eecs.berkeley.edu/~trevor/CS280Notes/>)
- Bill Freeman
- Kristen Grauman
- Tinne Tuytelaars
- And many others...

Computer Vision 1

(total #slides 169 | Lecture 4)

Summary

- 1. Feature Detection**
- 2. Motion**
- 3. Classification -> next week!**